

Motivation

- Leaky Integrate-and-Fire Neurons
 - work by integrating (accumulating) input currents over time to increase a "membrane potential" and "fire" (produce a spike) when this potential crosses a threshold, similar to how a leaky RC circuit charges. After firing, the potential resets. The "leaky" part refers to the fact that the membrane potential also gradually decays or "leaks" back to a resting value over time unless it is being driven by input
- The SCN idea
 - a population of LIF neurons only fire when it reduces the global representation error. This yields sparse, irregular, yet coordinated spikes
 - The paper shows that the same network can implement a linear dynamical system (add a "slow" recurrent term).
 - Kalman filter for optimal state estimation through noisy observations
 - LQR for optimal control to a target $z(t)$
 - The control u is a linear readout of the spikes
 - -> Spiking LQG

Setup




Fig. 1: Controlling dynamical systems with spiking neural networks. A recurrent spiking neural network (left) will emit spikes based on its inputs and connectivity, which are translated into a control signal to control some target system (right). F_c and F_m represent input weights, Ω represent recurrent weights, and D represent read-out weights. By first defining the optimal control solution for the system, and then directly translating the resulting control parameters into these connection weights, we can generate a robust and population-wide coordinated spiking network that accurately controls identified systems.

- x : state
- y : incomplete measurement of the system state
- z : incomplete measurement of the target state
- u : control signal
- s : spiking signal emitted from a recurrent SNN, used to control the system: filtered spikes

Emitted Spikes and Filtered Spikes

- s : spiking signal emitted from a recurrent SNN, used to control the state

$$s_i(t) = \sum_k \delta(t - t_i^{(k)}), \quad s(t) = [s_1(t), \dots, s_N(t)]^\top.$$

○

$$s_i[n] = \begin{cases} \frac{1}{\Delta t}, & \text{if neuron } i \text{ spikes in bin } n, \\ 0, & \text{otherwise,} \end{cases}$$

○ discrete:

○ $s_i(t)$ is the sum of the spikes emitted from the neurons in time t .

• Each $t_i^{(k)}$ is the k -th spike time of neuron i .

• Integrating s_i over any interval counts spikes: $\int_{t_0}^{t_1} s_i(t) dt = \#\{\text{spikes of } i \text{ in } [t_0, t_1]\}$.

• This is a **distribution** (generalized function), not an ordinary continuous signal. In their LIF network, a spike occurs when the voltage $v_i(t)$ hits its threshold T_i ; then s_i contributes a delta at that instant.

○

- r : filtered spikes

$$\dot{r}(t) = -\lambda r(t) + s(t) \Rightarrow r = (e^{-\lambda t} H(t)) * s,$$

○ so each spike adds a unit jump to r that then decays with time constant $1/\lambda$.

Closed-form control with spike ...

$$\dot{r}(t) = -\lambda r(t) + s(t) \iff r(t) = \int_{-\infty}^t e^{-\lambda(t-\tau)} s(\tau) d\tau.$$

○

- Dirac impulse delta function

$$\delta(x) = \begin{cases} 0, & x \neq 0 \\ \infty, & x = 0 \end{cases} \quad \text{with} \quad \int_{-\infty}^{\infty} \delta(x) dx = 1.$$

○

○ it's not a real function, so approximated with limits: $\delta_a(x) = \frac{1}{|a| \sqrt{\pi}} e^{-(x/a)^2}$ $a \rightarrow 0$

Decoding matrix

- Assumption

○ the input signal estimate $\hat{x}$ can be linearly decoded from the spiking activity as

$\hat{x} = Dr$, where D in $\mathbb{R}^{(N \times K)}$ are known decoding weights (in this paper randomly drawn from the normal distribution and normalized)

○ spikes should only be emitted when this improves the coding error defined by the L2-norm (yielding a GREEDY SPIKING RULE): neuron i should spike iff

$$\|x - Dr\|_2^2 > \|x - Dr - D_i\|_2^2$$

- Definition

$$\hat{x}(t) \in \mathbb{R}^K = D r(t), \quad r(t) \in \mathbb{R}^N.$$

So each **column** $D_i \in \mathbb{R}^K$ says “how a spike from neuron i moves the estimate.” Because a spike makes r_i jump by 1, it adds exactly D_i to $\hat{x}$. The paper also uses $D^\top$ as the **forward (encoding) weights** that

- inject the input x into voltages, which fixes the orientation: D is $K \times N$ (so $D^\top$ is $N \times K$).

In short: **D is the $K \times N$ dictionary that decodes the network’s filtered spikes into the state estimate; its columns are the one-spike “basis vectors,” and it induces both thresholds and the fast**

- **coordination term** $-D^\top D$.

- How to choose D

- sample the columns of D randomly (Gaussian) and normalize them

If you’re building this:

- - pick $N \geq K$ (often $N \gg K$),
 - draw $D_i \sim \mathcal{N}(0, I_K)$, then normalize each D_i (column),
 - keep this D fixed; it defines the coding dictionary and all induced weights.

- Calculating Threshold

- $\|e\|_2^2 > \|e - D_i\|_2^2$
- $\|e\|_2^2 > \|e\|_2^2 - 2D_i^\top e + \|D_i\|_2^2$
- $0 > 0 - 2D_i^\top e + \|D_i\|_2^2$
- $D_i^\top e > \frac{1}{2}\|D_i\|_2^2$
- $v_i > \frac{1}{2}\|D_i\|_2^2$
- so set the threshold T_i as such

Fast Recurrent Connection: Omega_f

- Target behavior we want

- voltages should always equal “projection of the current error onto each decoder”
- $v_j = D_j^\top (x - \hat{x}) = D_j^\top e$

- What happens after neuron i spikes

- $\hat{x}$ becomes $\hat{x}' = \hat{x} + D_i$
- e becomes $e' = x - \hat{x}' = e - D_i$
- the new correct voltage v_j becomes

$$v_j' = D_j^\top e' = D_j^\top e - D_j^\top D_i = v_j - D_j^\top D_i$$

- How to implement that instantly: choose the fast recurrent weight matrix
 - $\Omega_f = -D^T D$
 - No double-encoding: neurons whose decoders look like D_i (large $D_j^T D_i$) get their voltages pulled down the most, so they don't reflexively fire to add the same direction again.
 - Self-reset: the diagonal entry $-||D||^2$ knocks down v_i strongly after it spikes, preventing immediate re-firing unless the residual still demands it.

Voltage dynamics for tracking a fully observed x

- $\dot{v} = -\lambda v + D^T(\dot{x} + \lambda x) + \Omega_f s$
- Proof
 - $\dot{v} = D^T(\dot{x} - \dot{\hat{x}})$
 - $\dot{v} = D^T(\dot{x} - D\dot{r})$
 - $\dot{v} = D^T\dot{x} - D^T D\dot{r}$
 - $\dot{v} = D^T\dot{x} - D^T D(-\lambda r + s)$
 - $\dot{v} = D^T\dot{x} + D^T D\lambda r - D^T Ds$
 - notice $v = D^T x - D^T D r$
 - giving $D^T D r = D^T x - v$
 - $\dot{v} = -\lambda v + D^T(\dot{x} + \lambda x) - D^T Ds$
 - $\dot{v} = -\lambda v + D^T(\dot{x} + \lambda x) + \Omega_f s$

Spike Coding Theory

- If a K-dimensional state $x(t) \in \mathbb{R}^K$ is fully observable without noise, SCN theory defines how a recurrent SNN can optimally track this signal
 - Assumptions
 - the input signal estimate $\hat{x}$ can be linearly decoded from the spiking activity as $\hat{x} = D r$, where D in $\mathbb{R}^{(N \times K)}$ are known decoding weights (in this paper randomly drawn from the normal distribution and normalized)
 - spikes should only be emitted when this improves the coding error defined by the L2-norm (yielding a GREEDY SPIKING RULE): neuron i should spike iff $||x - D r||_2^2 > ||x - D r - D_i||_2^2$
- The main equations

- $\dot{\hat{x}} = Dr$
- $\dot{r} = -\lambda r + s$
- $\Omega_f = -D^T D$
- $\dot{v} = -\lambda v + D^T(\dot{x} + \lambda x) + \Omega_f s$